

## ▼ Preprocessing

```
# Imports
import pandas as pd
import numpy as np
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, GlobalAveragePooling1D, Dense
from tensorflow.keras.layers import SimpleRNN
from tensorflow.keras.optimizers import Adam
```

Double-click (or enter) to edit

```
# Load sms-spam data
data = pd.read_csv('sample_data/spam.csv', encoding = 'latin-1')
data = data.iloc[:, :2]
data.columns = ['label', 'text']

# Preprocessing
data['label'] = data['label'].apply(lambda x: 1 if x == 'spam' else 0)

# Split into 2 sets: Count and tf-idf
train_data, test_data = train_test_split(data, test_size=0.2, random_state=42)
```

## ▼ Distribution and Description

```
# Plot the distribution of the target classes
plt.hist(data['label'], bins=2, width=.2, align='right')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks([.35, .85], ['ham', 'spam'])
plt.show()
```



The [dataset](#) contains over 5000 SMS messages which are either spam or "ham" as shown in the distribution above. Messages that are "ham" are simply not spam. The model should be able to predict which messages are spam and which are "ham" or not spam based on the message. It should be able to analyze frequent words that appear across spam messages and use that to identify new spam messages.

## ▼ Sequential Model

```
# Tokenize and pad the text
tokenizer = Tokenizer()
tokenizer.fit_on_texts(train_data['text'])
vocab_size = len(tokenizer.word_index) + 1

# Train and test
X_train = tokenizer.texts_to_sequences(train_data['text'])
X_test = tokenizer.texts_to_sequences(test_data['text'])

max_length = max(len(x) for x in X_train)
X_train = pad_sequences(X_train, maxlen=max_length, padding='post')
X_test = pad_sequences(X_test, maxlen=max_length, padding='post')

y_train = train_data['label'].values
y_test = test_data['label'].values

# Define the sequential model
model = Sequential([
    Embedding(vocab_size, 16, input_length=max_length),
    GlobalAveragePooling1D(),
    Dense(16, activation='relu'),
    Dense(1, activation='sigmoid')
])

# Compile and train the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model.fit(X_train, y_train, epochs=20, validation_split=0.1)

# Evaluate the model on the test data
loss, accuracy = model.evaluate(X_test, y_test)
print(f"Test accuracy: {accuracy}")

Epoch 1/20
126/126 [=====] - 2s 5ms/step - loss: 0.5303 - accuracy: 0.8659 - val_loss: 0.3954
Epoch 2/20
126/126 [=====] - 0s 4ms/step - loss: 0.3580 - accuracy: 0.8671 - val_loss: 0.3703
Epoch 3/20
126/126 [=====] - 1s 4ms/step - loss: 0.3354 - accuracy: 0.8671 - val_loss: 0.3445
Epoch 4/20
126/126 [=====] - 0s 4ms/step - loss: 0.2971 - accuracy: 0.8671 - val_loss: 0.2936
Epoch 5/20
126/126 [=====] - 1s 4ms/step - loss: 0.2265 - accuracy: 0.8838 - val_loss: 0.2105
Epoch 6/20
126/126 [=====] - 0s 4ms/step - loss: 0.1451 - accuracy: 0.9484 - val_loss: 0.1441
Epoch 7/20
126/126 [=====] - 1s 4ms/step - loss: 0.0904 - accuracy: 0.9766 - val_loss: 0.1098
Epoch 8/20
```

```

126/126 [=====] - 0s 4ms/step - loss: 0.0615 - accuracy: 0.9838 - val_loss: 0.0905
Epoch 9/20
126/126 [=====] - 1s 5ms/step - loss: 0.0467 - accuracy: 0.9880 - val_loss: 0.0806
Epoch 10/20
126/126 [=====] - 0s 4ms/step - loss: 0.0375 - accuracy: 0.9900 - val_loss: 0.0772
Epoch 11/20
126/126 [=====] - 1s 4ms/step - loss: 0.0312 - accuracy: 0.9910 - val_loss: 0.0731
Epoch 12/20
126/126 [=====] - 1s 4ms/step - loss: 0.0265 - accuracy: 0.9925 - val_loss: 0.0745
Epoch 13/20
126/126 [=====] - 0s 4ms/step - loss: 0.0233 - accuracy: 0.9935 - val_loss: 0.0645
Epoch 14/20
126/126 [=====] - 1s 4ms/step - loss: 0.0203 - accuracy: 0.9935 - val_loss: 0.0690
Epoch 15/20
126/126 [=====] - 1s 4ms/step - loss: 0.0177 - accuracy: 0.9953 - val_loss: 0.0661
Epoch 16/20
126/126 [=====] - 0s 4ms/step - loss: 0.0158 - accuracy: 0.9958 - val_loss: 0.0665
Epoch 17/20
126/126 [=====] - 1s 4ms/step - loss: 0.0141 - accuracy: 0.9963 - val_loss: 0.0651
Epoch 18/20
126/126 [=====] - 1s 4ms/step - loss: 0.0124 - accuracy: 0.9973 - val_loss: 0.0686
Epoch 19/20
126/126 [=====] - 1s 5ms/step - loss: 0.0116 - accuracy: 0.9965 - val_loss: 0.0665
Epoch 20/20
126/126 [=====] - 1s 6ms/step - loss: 0.0102 - accuracy: 0.9968 - val_loss: 0.0687
35/35 [=====] - 0s 2ms/step - loss: 0.0678 - accuracy: 0.9821
Test accuracy: 0.9820627570152283

```

## ▼ RNN

```

# Define the RNN model
model_rnn = Sequential([
    Embedding(vocab_size, 16, input_length=max_length),
    SimpleRNN(32),
    Dense(1, activation='sigmoid')
])

# Modify learning rate
optimizer = Adam(learning_rate=.0001)

# Compile and train the model
model_rnn.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accuracy'])
model_rnn.fit(X_train, y_train, epochs=20, validation_split=0.1)

# Evaluate the model on the test data
loss, accuracy = model_rnn.evaluate(X_test, y_test)
print(f"Test accuracy (RNN): {accuracy}")

Epoch 1/20
126/126 [=====] - 6s 32ms/step - loss: 0.5063 - accuracy: 0.8340 - val_loss: 0.4231
Epoch 2/20
126/126 [=====] - 4s 30ms/step - loss: 0.3983 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 3/20
126/126 [=====] - 5s 43ms/step - loss: 0.3925 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 4/20
126/126 [=====] - 4s 30ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 5/20
126/126 [=====] - 4s 29ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 6/20
126/126 [=====] - 5s 41ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 7/20
126/126 [=====] - 4s 29ms/step - loss: 0.3919 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 8/20
126/126 [=====] - 4s 30ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 9/20
126/126 [=====] - 5s 41ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.4111
Epoch 10/20

```

```

126/126 [=====] - 4s 30ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.411
Epoch 11/20
126/126 [=====] - 4s 31ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.411
Epoch 12/20
126/126 [=====] - 5s 41ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.411
Epoch 13/20
126/126 [=====] - 4s 30ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.411
Epoch 14/20
126/126 [=====] - 4s 31ms/step - loss: 0.3921 - accuracy: 0.8671 - val_loss: 0.411
Epoch 15/20
126/126 [=====] - 6s 44ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.411
Epoch 16/20
126/126 [=====] - 6s 45ms/step - loss: 0.3919 - accuracy: 0.8671 - val_loss: 0.411
Epoch 17/20
126/126 [=====] - 4s 30ms/step - loss: 0.3922 - accuracy: 0.8671 - val_loss: 0.411
Epoch 18/20
126/126 [=====] - 5s 41ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.411
Epoch 19/20
126/126 [=====] - 4s 30ms/step - loss: 0.3920 - accuracy: 0.8671 - val_loss: 0.411
Epoch 20/20
126/126 [=====] - 4s 30ms/step - loss: 0.3921 - accuracy: 0.8671 - val_loss: 0.411
35/35 [=====] - 0s 8ms/step - loss: 0.3949 - accuracy: 0.8655
Test accuracy (RNN): 0.865470826625824

```

## ▼ GloVe embedding

```

# Load 100D GloVe embeddings
glove_file_path = 'sample_data/glove.6B.100d.txt'
embeddings_index = {}
with open(glove_file_path, encoding='utf-8') as file:
    for line in file:
        values = line.split()
        word = values[0]
        coefs = np.asarray(values[1:], dtype='float32')
        embeddings_index[word] = coefs

# Create the embedding matrix
embedding_dim = 100
embedding_matrix = np.zeros((vocab_size, embedding_dim))
for word, i in tokenizer.word_index.items():
    if i < vocab_size:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector

# Define the model with GloVe embeddings
model_glove = Sequential([
    Embedding(vocab_size, embedding_dim, input_length=max_length, weights=[embedding_matrix], trainable=False),
    SimpleRNN(32),
    Dense(1, activation='sigmoid')
])

# Compile and train the model
model_glove.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
model_glove.fit(X_train, y_train, epochs=20, validation_split=0.1)

# Evaluate the model on the test data
loss, accuracy = model_glove.evaluate(X_test, y_test)
print(f"Test accuracy (GloVe): {accuracy}")

Epoch 1/20
126/126 [=====] - 5s 28ms/step - loss: 0.4335 - accuracy: 0.8070 - val_loss: 0.360
Epoch 2/20
126/126 [=====] - 5s 38ms/step - loss: 0.3483 - accuracy: 0.8671 - val_loss: 0.350

```

```

Epoch 3/20
126/126 [=====] - 3s 25ms/step - loss: 0.3214 - accuracy: 0.8621 - val_loss: 0.354
Epoch 4/20
126/126 [=====] - 3s 25ms/step - loss: 0.3053 - accuracy: 0.8659 - val_loss: 0.440
Epoch 5/20
126/126 [=====] - 3s 25ms/step - loss: 0.3933 - accuracy: 0.8669 - val_loss: 0.417
Epoch 6/20
126/126 [=====] - 5s 37ms/step - loss: 0.3900 - accuracy: 0.8669 - val_loss: 0.401
Epoch 7/20
126/126 [=====] - 3s 25ms/step - loss: 0.3849 - accuracy: 0.8669 - val_loss: 0.399
Epoch 8/20
126/126 [=====] - 3s 25ms/step - loss: 0.3637 - accuracy: 0.8646 - val_loss: 0.330
Epoch 9/20
126/126 [=====] - 4s 29ms/step - loss: 0.3687 - accuracy: 0.8549 - val_loss: 0.411
Epoch 10/20
126/126 [=====] - 4s 33ms/step - loss: 0.3930 - accuracy: 0.8669 - val_loss: 0.411
Epoch 11/20
126/126 [=====] - 3s 25ms/step - loss: 0.3931 - accuracy: 0.8669 - val_loss: 0.411
Epoch 12/20
126/126 [=====] - 3s 25ms/step - loss: 0.3923 - accuracy: 0.8669 - val_loss: 0.411
Epoch 13/20
126/126 [=====] - 4s 36ms/step - loss: 0.3925 - accuracy: 0.8671 - val_loss: 0.411
Epoch 14/20
126/126 [=====] - 3s 25ms/step - loss: 0.3930 - accuracy: 0.8671 - val_loss: 0.411
Epoch 15/20
126/126 [=====] - 3s 25ms/step - loss: 0.3925 - accuracy: 0.8671 - val_loss: 0.411
Epoch 16/20
126/126 [=====] - 3s 25ms/step - loss: 0.3935 - accuracy: 0.8671 - val_loss: 0.411
Epoch 17/20
126/126 [=====] - 5s 37ms/step - loss: 0.3933 - accuracy: 0.8671 - val_loss: 0.411
Epoch 18/20
126/126 [=====] - 3s 26ms/step - loss: 0.3927 - accuracy: 0.8671 - val_loss: 0.411
Epoch 19/20
126/126 [=====] - 3s 25ms/step - loss: 0.3930 - accuracy: 0.8671 - val_loss: 0.411
Epoch 20/20
126/126 [=====] - 4s 28ms/step - loss: 0.3917 - accuracy: 0.8671 - val_loss: 0.411
35/35 [=====] - 1s 14ms/step - loss: 0.4009 - accuracy: 0.8637
Test accuracy (GloVe): 0.8636771440505981

```

## ▼ Results

```

# Sequential
_, seq_accuracy = model.evaluate(X_test, y_test, verbose=0)

# SimpleRNN
_, rnn_accuracy = model_rnn.evaluate(X_test, y_test, verbose=0)

# SimpleRNN with GloVe
_, glove_accuracy = model_glove.evaluate(X_test, y_test, verbose=0)

# Print the accuracy scores for each model
print(f"Test accuracy (Sequential): {seq_accuracy:.2f}")
print(f"Test accuracy (SimpleRNN): {rnn_accuracy:.2f}")
print(f"Test accuracy (SimpleRNN with GloVe embeddings): {glove_accuracy:.2f}")

Test accuracy (Sequential): 0.98
Test accuracy (SimpleRNN): 0.87
Test accuracy (SimpleRNN with GloVe embeddings): 0.86

```

## Summary & Analysis

The highest performing deep learning model was actually the sequential model, which is the most simple implementation. It makes sense that it works better because this dataset is very simple and a very simple instance of classification.

## RNN

The SimpleRNN capped at 86.7% with no special embeddings. This was lower than the sequential model, but it is still high and it is able to reach this accuracy after very few epochs.

## Embeddings

I implemented GloVe embeddings, a powerful word embedding model. It actually reduced the SimpleRNN's learning capabilities, making it take more epochs in order to reach the same max accuracy.

---

✓ 1s completed at 5:21 PM

